

Problem: Predicting Airplane Delays

The goals of this notebook are:

- Process and create a dataset from downloaded ZIP files
- Exploratory data analysis (EDA)
- Establish a baseline model and improve it

Introduction to business scenario

You work for a travel booking website that is working to improve the customer experience for flights that were delayed. The company wants to create a feature to let customers know if the flight will be delayed due to weather when the customers are booking the flight to or from the busiest airports for domestic travel in the US.

You are tasked with solving part of this problem by leveraging machine learning to identify whether the flight will be delayed due to weather. You have been given access to the a dataset of on-time performance of domestic flights operated by large air carriers. You can use this data to train a machine learning model to predict if the flight is going to be delayed for the busiest airports.

Dataset

The provided dataset contains scheduled and actual departure and arrival times reported by certified US air carriers that account for at least 1 percent of domestic scheduled passenger revenues. The data was collected by the Office of Airline Information, Bureau of Transportation Statistics (BTS). The dataset contains date, time, origin, destination, airline, distance, and delay status of flights for flights between 2014 and 2018. The data are in 60 compressed files, where each file contains a CSV for the flight details in a month for the five years (from 2014 - 2018). The data can be downloaded from this [link](https://ucstaff-my.sharepoint.com/:f/g/personal/ibrahim_radwan_camberra_edu_au/EhWeqeQsh-9Mr1fneZc9_0sBOBzEdXngvxFJtAlla-eAgA?e=8ukWwa). Please download the data files and place them on a relative path. Dataset(s) used in this assignment were compiled by the Office of Airline Information, Bureau of Transportation Statistics (BTS), Airline On-Time Performance Data, available with the following [link](https://www.transtats.bts.gov/Fields.asp?gnoyr_VQ=FGJ).

Step 1: Prepare the environment

Use one of the labs which we have practised on with the Amazon Sagemakers where you perform the following steps:

1. Start a lab.
2. Create a notebook instance and name it "oncloudproject".
3. Increase the used memory to 25 GB from the additional configurations.
4. Open Jupyter Lab and upload this notebook into it.
5. Upload the two combined CSV files (combined_csv_v1.csv and combined_csv_v2.csv), which you created in Part A of this project.

Note: In case of the data is too much to be uploaded to the AWS, please use 20% of the data only for this task.

Step 2: Build and evaluate simple models

Write code to perform the following steps:

1. Split data into training, validation and testing sets (70% - 15% - 15%).
2. Use linear learner estimator to build a classification model.
3. Host the model on another instance
4. Perform batch transform to evaluate the model on testing data
5. Report the performance metrics that you see better test the model performance

Note: You are required to perform the above steps on the two combined datasets separately and to comment on the difference.

```
In [1]: # Import libraries
import boto3, sagemaker, io, os, warnings
import pandas as pd
from sklearn.model_selection import train_test_split
warnings.simplefilter('ignore')

import sagemaker
import numpy as np
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt
from sagemaker import image_uris
from sagemaker.image_uris import retrieve
from sklearn.model_selection import train_test_split

sagemaker.config INFO - Not applying SDK defaults from location: /etc/xdg/sagemaker/config.yaml
sagemaker.config INFO - Not applying SDK defaults from location: /home/ec2-user/.config/sagemaker/config.yaml
```

```
In [2]: # Load dataset
df = pd.read_csv('combined_csv_v1.csv')

# sample 5% because data is large
df = df.sample(frac=0.05, random_state=42)

# Quick check
df.shape
```

Out[2]: (81780, 75)

```
In [3]: (df['target']==1).sum()
```

Out[3]: 17100

```
In [4]: (df['target']==0).sum()
```

Out[4]: 64680

```
In [5]: df.head()
```

	target	Distance	Quarter_2	Quarter_3	Quarter_4	Month_2	Month_3	Month_4	Month_5	Month_6	...	Dest_DEN	Dest_DFW	Dest_IAH	...
1312591	0.0	2125.0	True	False	False	False	False	False	True	False	...	False	False	False	...
7589	1.0	731.0	False	False	True	False	False	False	False	False	...	False	False	False	...
1438862	0.0	1587.0	False	False	True	False	False	False	False	False	...	False	False	False	...
434159	0.0	1464.0	True	False	False	False	False	False	True	False	...	False	True	False	...
48589	0.0	1464.0	False	False	False	True	False	False	False	False	...	False	True	False	...

5 rows × 75 columns

```
In [6]: pd.set_option('display.max_rows', None)
print(df.dtypes)
```

target	float64
Distance	float64
Quarter_2	bool
Quarter_3	bool
Quarter_4	bool
Month_2	bool
Month_3	bool
Month_4	bool
Month_5	bool
Month_6	bool
Month_7	bool
Month_8	bool
Month_9	bool
Month_10	bool
Month_11	bool
Month_12	bool
DayofMonth_2	bool
DayofMonth_3	bool
DayofMonth_4	bool
DayofMonth_5	bool

1. Split data into training, validation and testing sets (70% - 15% - 15%).

```
In [6]: # I have issue with batch transforming because features are not integer, so I convert all boolean columns to
df = df.astype({col: 'int64' for col in df.select_dtypes(include=['bool']).columns})

# Check the result:
print("Data types after conversion:")
print(df.dtypes.value_counts())

# Move target column to first position (required for Linear Learner)
cols = list(df.columns)
cols.remove('target')
cols = ['target'] + cols
df = df[cols]

print(f"\nFinal data shape: {df.shape}")
print(f"First column (should be target): {df.columns[0]}")
```

Data types after conversion:

int64 73

float64 2

Name: count, dtype: int64

Final data shape: (81780, 75)

First column (should be target): target

```
In [7]: pd.set_option('display.max_rows', None)
print(df.dtypes)
```

target	float64
Distance	float64
Quarter_2	int64
Quarter_3	int64
Quarter_4	int64
Month_2	int64
Month_3	int64
Month_4	int64
Month_5	int64
Month_6	int64
Month_7	int64
Month_8	int64
Month_9	int64
Month_10	int64
Month_11	int64
Month_12	int64
DayofMonth_2	int64
DayofMonth_3	int64
DayofMonth_4	int64
DayofMonth_5	int64

In [8]: # Split 70-15-15

```
train, test_validate = train_test_split(  
    df,  
    test_size=0.30, # 30% reserved for test + validation  
    random_state=42,  
    stratify=df['target'])  
  
validate, test = train_test_split(  
    test_validate,  
    test_size=0.50, # split the 30% into two halves = 15% each  
    random_state=42,  
    stratify=test_validate['target'])  
  
print("Train:", train.shape)  
print("Validate:", validate.shape)  
print("Test:", test.shape)  
  
print("\nTrain distribution:\n", train['target'].value_counts(normalize=True))  
print("\nValidate distribution:\n", validate['target'].value_counts(normalize=True))  
print("\nTest distribution:\n", test['target'].value_counts(normalize=True))
```

```
Train: (57246, 75)
Validate: (12267, 75)
Test: (12267, 75)

Train distribution:
target
0.0    0.790902
1.0    0.209098
Name: proportion, dtype: float64

Validate distribution:
target
0.0    0.790902
1.0    0.209098
Name: proportion, dtype: float64

Test distribution:
target
0.0    0.790902
1.0    0.209098
Name: proportion, dtype: float64
```

In []:

2. Use linear learner estimator to build a classification model.

Upload to S3

In [9]: `bucket='c182567a4701747l12448301t1w754121081532-labbucket-0p0w0femass0'`

```
In [10]: prefix='lab3'

train_file='delay_train.csv'
test_file='delay_test.csv'
validate_file='delay_validate.csv'

s3_resource = boto3.Session().resource('s3')
def upload_s3_csv(filename, folder, dataframe):
    csv_buffer = io.StringIO()
    dataframe.to_csv(csv_buffer, header=False, index=False)
    s3_resource.Bucket(bucket).Object(os.path.join(prefix, folder, filename)).put(Body=csv_buffer.getvalue())

In [11]: upload_s3_csv(train_file, 'train', train)
upload_s3_csv(test_file, 'test', test)
upload_s3_csv(validate_file, 'validate', validate)
```

```
In [12]: for obj in s3_resource.Bucket(bucket).objects.filter(Prefix=prefix):
    print(obj.key)
```

```
lab3/batch-in/batch-in.csv
lab3/output/linear-learner-2025-11-03-07-52-53-188/debug-output/training_job_end.ts
lab3/output/linear-learner-2025-11-03-07-52-53-188/output/model.tar.gz
lab3/output/linear-learner-2025-11-03-07-52-53-188/profiler-output/framework/training_job_end.ts
lab3/output/linear-learner-2025-11-03-07-52-53-188/profiler-output/system/incremental/2025110307/176215638
0.algo-1.json
lab3/output/linear-learner-2025-11-03-07-52-53-188/profiler-output/system/incremental/2025110307/176215644
0.algo-1.json
lab3/output/linear-learner-2025-11-03-07-52-53-188/profiler-output/system/incremental/2025110307/176215650
0.algo-1.json
lab3/output/linear-learner-2025-11-03-07-52-53-188/profiler-output/system/incremental/2025110307/176215656
0.algo-1.json
lab3/output/linear-learner-2025-11-03-07-52-53-188/profiler-output/system/incremental/2025110307/176215662
0.algo-1.json
lab3/output/linear-learner-2025-11-03-07-52-53-188/profiler-output/system/training_job_end.ts
lab3/output/linear-learner-2025-11-03-08-25-21-441/debug-output/training_job_end.ts
lab3/output/linear-learner-2025-11-03-08-25-21-441/output/model.tar.gz
lab3/output/linear-learner-2025-11-03-08-25-21-441/profiler-output/framework/training_job_end.ts
lab3/output/linear-learner-2025-11-03-08-25-21-441/profiler-output/system/incremental/2025110308/176215836
0.algo-1.json
lab3/output/linear-learner-2025-11-03-08-25-21-441/profiler-output/system/incremental/2025110308/176215842
0.algo-1.json
lab3/output/linear-learner-2025-11-03-08-25-21-441/profiler-output/system/incremental/2025110308/176215848
0.algo-1.json
lab3/output/linear-learner-2025-11-03-08-25-21-441/profiler-output/system/incremental/2025110308/176215854
0.algo-1.json
lab3/output/linear-learner-2025-11-03-08-25-21-441/profiler-output/system/incremental/2025110308/176215860
0.algo-1.json
lab3/output/linear-learner-2025-11-03-08-25-21-441/profiler-output/system/incremental/2025110308/176215866
0.algo-1.json
lab3/output/linear-learner-2025-11-03-08-25-21-441/profiler-output/system/incremental/2025110308/176215872
0.algo-1.json
lab3/output/linear-learner-2025-11-03-08-25-21-441/profiler-output/system/training_job_end.ts
lab3/test/delay_test.csv
lab3/train/delay_train.csv
lab3/validate/delay_validate.csv
```

Training Model

```
In [13]: container = image_uris.retrieve(framework='linear-learner', region=boto3.Session().region_name)
```

```
In [14]: # Set hyperparameters for Linear Learner
hyperparams = {
    'feature_dim': str(train.shape[1] - 1), # Number of features (exclude target)
    'predictor_type': 'binary_classifier',
    'mini_batch_size': '100',
    'epochs': '10',
    'num_models': '32',
    'loss': 'logistic',
    'learning_rate': 'auto',
    'normalize_data': 'true',
    'normalize_label': 'false',
    'unbias_data': 'false',
    'unbias_label': 'false',
    'num_calibration_samples': '10000000',
    'early_stopping_patience': '3',
    'early_stopping_tolerance': '0.001'
}
```

```
In [15]: # Create output path
s3_output_location = "s3://{} / {} / output / ".format(bucket, prefix)
```

```
In [16]: # Create Linear Learner estimator
linear_model = sagemaker.estimator.Estimator(
    container,
    sagemaker.get_execution_role(),
    instance_count=1,
    instance_type='ml.m4.xlarge',
    output_path=s3_output_location,
    hyperparameters=hyperparams,
    sagemaker_session=sagemaker.Session()
)
```

```
In [17]: # Set up data channels
train_channel = sagemaker.inputs.TrainingInput(
    's3://{}:{}/train/{}'.format(bucket, prefix, train_file),
    content_type='text/csv'
)
validate_channel = sagemaker.inputs.TrainingInput(
    's3://{}:{}/validate/{}'.format(bucket, prefix, validate_file),
    content_type='text/csv'
)
```

```
In [18]: # Train the model
linear_model.fit(
    inputs={
        'train': train_channel,
        'validation': validate_channel
    },
    wait=True
)
print("Training completed!")
```

```
INFO:sagemaker:Creating training-job with name: linear-learner-2025-11-03-08-37-28-126
2025-11-03 08:37:29 Starting - Starting the training job...
2025-11-03 08:37:53 Starting - Preparing the instances for training...
2025-11-03 08:38:23 Downloading - Downloading input data...
2025-11-03 08:38:53 Downloading - Downloading the training image.....
2025-11-03 08:40:24 Training - Training image download completed. Training in progress
.
.
.
2025-11-03 08:45:33 Uploading - Uploading generated training model
2025-11-03 08:45:33 Completed - Training job completed
..Training seconds: 430
Billable seconds: 430
Training completed!
```

3. Host the model on another instance

```
In [19]: # Deploy the trained Linear Learner model
linear_predictor = linear_model.deploy(
    initial_instance_count=1,
    instance_type='ml.m4.xlarge', # Can use different instance than training
)

print(f"Model deployed successfully!")
```

```
INFO:sagemaker:Creating model with name: linear-learner-2025-11-03-08-45-47-044
INFO:sagemaker:Creating endpoint-config with name linear-learner-2025-11-03-08-45-47-044
INFO:sagemaker:Creating endpoint with name linear-learner-2025-11-03-08-45-47-044
-----!Model deployed successfully!
```

4. Perform batch transform to evaluate the model on testing data

```
In [20]: # Prepare test data (features only, no target)
batch_X = test.iloc[:, 1:] # Exclude target column (first column)

# Upload batch input to S3
batch_X_file = 'batch-in.csv'
upload_s3_csv(batch_X_file, 'batch-in', batch_X)

# Define S3 paths
batch_output = "s3://{}{}/batch-out/".format(bucket, prefix)
batch_input = "s3://{}{}/batch-in/{}".format(bucket, prefix, batch_X_file)

print("Starting batch transform...")

# Create transformer
linear_transformer = linear_model.transformer(
    instance_count=1,
    instance_type='ml.m4.xlarge',
    strategy='MultiRecord',
    assemble_with='Line',
    output_path=batch_output
)

# Run transform
linear_transformer.transform(
    data=batch_input,
    data_type='S3Prefix',
    content_type='text/csv',
    split_type='Line'
)

# Wait for completion
linear_transformer.wait()
print(" Batch transform completed!")

# DOWNLOAD AND PARSE PREDICTIONS

import json

# Download predictions from S3
s3 = boto3.client('s3')
```

```
obj = s3.get_object(Bucket=bucket, Key="{}/batch-out/{}".format(prefix, 'batch-in.csv.out'))  
  
# Parse Linear Learner predictions (JSON format)  
predictions_raw = obj['Body'].read().decode('utf-8').strip().split('\n')  
  
# Extract predicted labels and scores  
predicted_labels = []  
predicted_scores = []  
  
for line in predictions_raw:  
    pred_json = json.loads(line)  
    predicted_labels.append(pred_json['predicted_label'])  
    predicted_scores.append(pred_json['score'])  
  
# Create DataFrame with predictions  
target_predicted = pd.DataFrame({  
    'target': predicted_labels,  
    'score': predicted_scores  
})
```

INFO:sagemaker:Creating model with name: linear-learner-2025-11-03-08-51-36-514

Starting batch transform...

INFO:sagemaker:Creating transform job with name: linear-learner-2025-11-03-08-51-37-082

.....
... Batch transform completed!

```
In [21]: # explore the prediction results
```

```
def binary_convert(x):
    threshold = 0.3
    if x > threshold:
        return 1
    else:
        return 0

target_predicted_binary = target_predicted['score'].apply(binary_convert)

print("Prediction DataFrame structure:")
print(target_predicted.head(10))
print("\nColumns:", target_predicted.columns.tolist())
print("\nData types:")
print(target_predicted.dtypes)
```

Prediction DataFrame structure:

	target	score
0	0	0.268004
1	0	0.080399
2	0	0.326775
3	0	0.065700
4	0	0.276970
5	0	0.436232
6	0	0.042667
7	0	0.192169
8	0	0.133488
9	0	0.169756

Columns: ['target', 'score']

Data types:

target	int64
score	float64
dtype:	object

```
In [25]: from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.metrics import confusion_matrix, classification_report, roc_auc_score

# Get true labels and predictions
y_true = test['target'].values
y_pred = target_predicted_binary.values
y_scores = target_predicted['score'].values

print("LINEAR LEARNER – PERFORMANCE METRICS (combined_csv_v1.csv)")

# Calculate all metrics
accuracy = accuracy_score(y_true, y_pred)
precision = precision_score(y_true, y_pred)
recall = recall_score(y_true, y_pred)
f1 = f1_score(y_true, y_pred)
roc_auc = roc_auc_score(y_true, y_scores)

print(f"\n1. Accuracy: {accuracy:.4f}")
print(f"2. Precision: {precision:.4f}")
print(f"3. Recall: {recall:.4f}")
print(f"4. F1-Score: {f1:.4f}")
print(f"5. ROC-AUC: {roc_auc:.4f}")

# Confusion Matrix
print("\n Confusion Matrix:")
cm = confusion_matrix(y_true, y_pred)
print(cm)
print(f"\n    True Negatives (TN): {cm[0,0]}")
print(f"    False Positives (FP): {cm[0,1]}")
print(f"    False Negatives (FN): {cm[1,0]}")
print(f"    True Positives (TP): {cm[1,1]}")
```

LINEAR LEARNER – PERFORMANCE METRICS (combined_csv_v1.csv)

1. Accuracy: 0.6856
2. Precision: 0.2711
3. Recall: 0.2982
4. F1-Score: 0.2840
5. ROC-AUC: 0.5765

Confusion Matrix:

```
[[7645 2057]
 [1800  765]]
```

True Negatives (TN): 7,645
False Positives (FP): 2,057
False Negatives (FN): 1,800
True Positives (TP): 765

```
In [26]: import matplotlib.pyplot as plt
import seaborn as sns

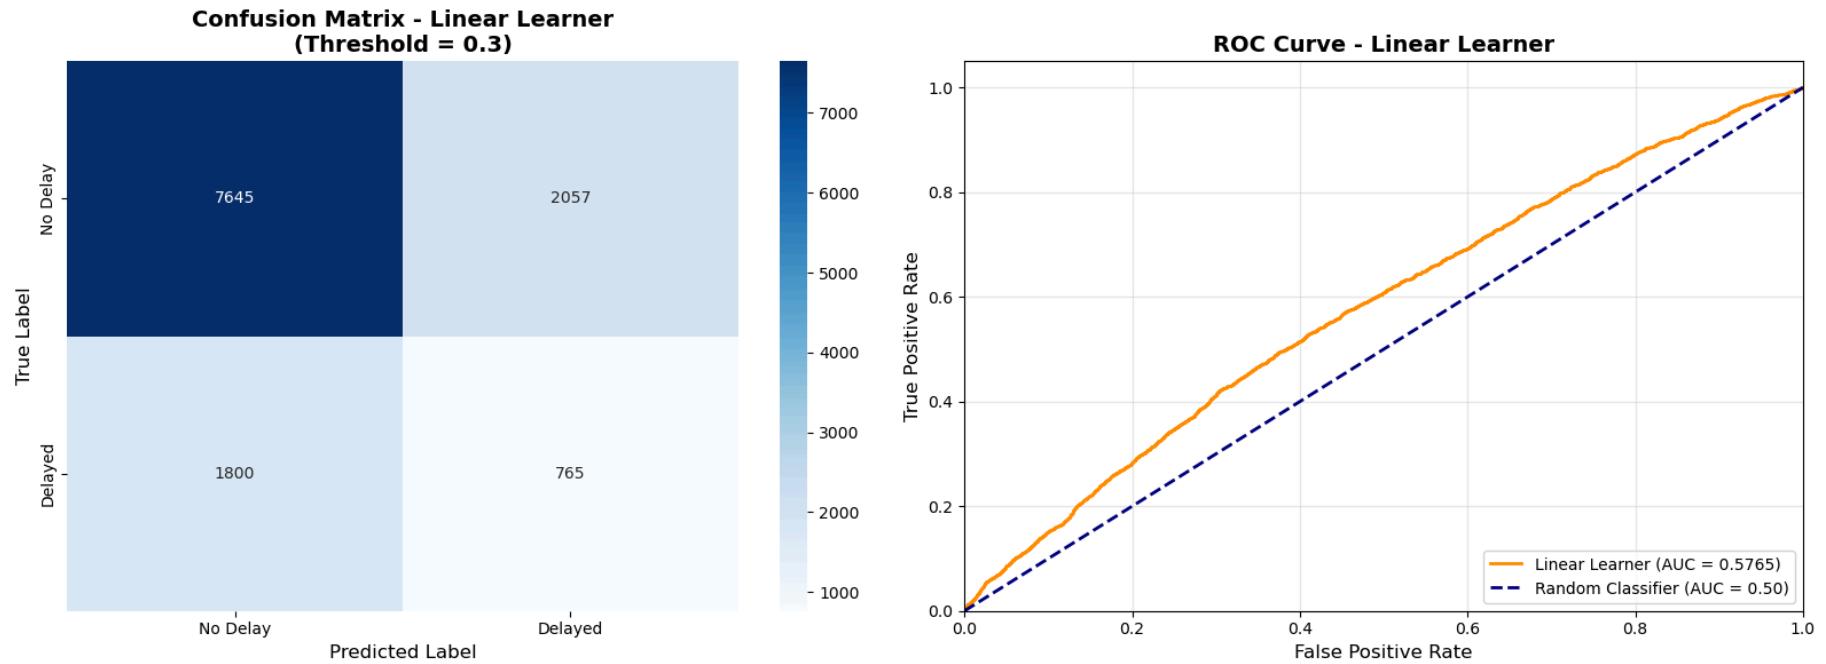
# Create figure with subplots
fig, axes = plt.subplots(1, 2, figsize=(16, 6))

# 1. Confusion Matrix Heatmap
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', ax=axes[0],
            xticklabels=['No Delay', 'Delayed'],
            yticklabels=['No Delay', 'Delayed'])
axes[0].set_title('Confusion Matrix - Linear Learner\nThreshold = 0.3', fontsize=14, fontweight='bold')
axes[0].set_ylabel('True Label', fontsize=12)
axes[0].set_xlabel('Predicted Label', fontsize=12)

# 2. ROC Curve
fpr, tpr, thresholds = roc_curve(y_true, y_scores)
roc_auc_calc = auc(fpr, tpr)

axes[1].plot(fpr, tpr, color='darkorange', lw=2,
             label=f'Linear Learner (AUC = {roc_auc_calc:.4f})')
axes[1].plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--',
             label='Random Classifier (AUC = 0.50)')
axes[1].set_xlim([0.0, 1.0])
axes[1].set_ylim([0.0, 1.05])
axes[1].set_xlabel('False Positive Rate', fontsize=12)
axes[1].set_ylabel('True Positive Rate', fontsize=12)
axes[1].set_title('ROC Curve - Linear Learner', fontsize=14, fontweight='bold')
axes[1].legend(loc="lower right", fontsize=10)
axes[1].grid(alpha=0.3)

plt.tight_layout()
plt.show()
```



Step 3: Build and evaluate ensemble models

Write code to perform the following steps:

1. Split data into training, validation and testing sets (70% - 15% - 15%).
2. Use xgboost estimator to build a classification model.
3. Host the model on another instance
4. Perform batch transform to evaluate the model on testing data
5. Report the performance metrics that you see better test the model performance
6. write down your observation on the difference between the performance of using the simple and ensemble models. Note: You are required to perform the above steps on the two combined datasets separately.

After base model, I use a new data called combined_csv_v2. It has new features added weather and holidays to test if there are any changes from the base model.

```
In [1]: # Import libraries
import boto3, sagemaker, io, os, warnings
import pandas as pd
from sklearn.model_selection import train_test_split
warnings.simplefilter('ignore')

import sagemaker
import numpy as np
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt
from sagemaker import image_uris
from sagemaker.image_uris import retrieve
from sklearn.model_selection import train_test_split

sagemaker.config INFO - Not applying SDK defaults from location: /etc/xdg/sagemaker/config.yaml
sagemaker.config INFO - Not applying SDK defaults from location: /home/ec2-user/.config/sagemaker/config.yaml
ml

Matplotlib is building the font cache; this may take a moment.
```

```
In [3]: # Load dataset
df = pd.read_csv('combined_csv_v2.csv')

# sample 5% because data is large
df = df.sample(frac=0.05, random_state=42)

# Quick check
df.shape
```

Out[3]: (3447, 86)

```
In [4]: (df['target']==1).sum()
```

Out[4]: 686

```
In [5]: (df['target']==0).sum()
```

Out[5]: 2761

In [6]: df.head()

Out[6]:

	target	Distance	DepHourOfDay	AWND_O	PRCP_O	TAVG_O	AWND_D	PRCP_D	TAVG_D	SNOW_O	...	Origin_SFO	Dest_CLT	Dest_DEN
68505	0.0	912.0	12	30	0	248.0	51	81	238.0	0.0	...	False	True	False
40439	0.0	602.0	13	66	0	82.0	23	0	166.0	0.0	...	False	False	False
63283	0.0	802.0	22	53	30	84.0	32	0	213.0	0.0	...	False	False	False
24901	0.0	606.0	11	47	0	-72.0	53	3	98.0	0.0	...	False	False	False
31780	0.0	1440.0	17	50	0	7.0	46	0	190.0	0.0	...	False	False	False

5 rows × 86 columns

```
In [7]: pd.set_option('display.max_rows', None)
print(df.dtypes)
```

target	float64
Distance	float64
DepHourOfDay	int64
AWND_0	int64
PRCP_0	int64
TAVG_0	float64
AWND_D	int64
PRCP_D	int64
TAVG_D	float64
SNOW_0	float64
SNOW_D	float64
Year_2015	bool
Year_2016	bool
Year_2017	bool
Year_2018	bool
Quarter_2	bool
Quarter_3	bool
Quarter_4	bool
Month_2	bool
Month_3	bool
Month_4	bool
Month_5	bool
Month_6	bool
Month_7	bool
Month_8	bool
Month_9	bool
Month_10	bool
Month_11	bool
Month_12	bool
DayofMonth_2	bool
DayofMonth_3	bool
DayofMonth_4	bool
DayofMonth_5	bool
DayofMonth_6	bool
DayofMonth_7	bool
DayofMonth_8	bool
DayofMonth_9	bool
DayofMonth_10	bool
DayofMonth_11	bool
DayofMonth_12	bool
DayofMonth_13	bool

DayofMonth_14	object
DayofMonth_15	object
DayofMonth_16	object
DayofMonth_17	object
DayofMonth_18	object
DayofMonth_19	object
DayofMonth_20	object
DayofMonth_21	object
DayofMonth_22	object
DayofMonth_23	object
DayofMonth_24	object
DayofMonth_25	object
DayofMonth_26	object
DayofMonth_27	object
DayofMonth_28	object
DayofMonth_29	object
DayofMonth_30	object
DayofMonth_31	object
DayOfWeek_2	object
DayOfWeek_3	object
DayOfWeek_4	object
DayOfWeek_5	object
DayOfWeek_6	object
DayOfWeek_7	object
Reporting_Airline_DL	object
Reporting_Airline_00	object
Reporting_Airline_UA	object
Reporting_Airline_WN	object
Origin_CLT	object
Origin_DEN	object
Origin_DFW	object
Origin_IAH	object
Origin_LAX	object
Origin_ORD	object
Origin_PHX	object
Origin_SF0	object
Dest_CLT	object
Dest_DEN	object
Dest_DFW	object
Dest_IAH	object
Dest_LAX	object
Dest_ORD	object

```
Dest_PHX          object
Dest_SF0          object
is_holiday_1     object
dtype: object
```

1. Split data into training, validation and testing sets (70% - 15% - 15%).

```
In [8]: # Convert all objects and boolean columns to 0/1

# df['Dest_IAH'].unique()
# df['is_holiday_1'].unique()

# First, I convert object columns that contain 'True'/'False' strings to real booleans
df = df.replace({'TRUE': True, 'FALSE': False, 'True': True, 'False': False})
```

```
In [9]: # df.dtypes
df = df.astype({col: 'int64' for col in df.select_dtypes(include=['bool']).columns})
# df.dtypes
```

```
In [10]: # Move target column to first position (required for Linear Learner)
cols = list(df.columns)
cols.remove('target')
cols = ['target'] + cols
df = df[cols]

print(f"\nFinal data shape: {df.shape}")
print(f"First column (should be target): {df.columns[0]}")
```

```
Final data shape: (3447, 86)
First column (should be target): target
```

In [11]: # Split 70-15-15

```
train, test_validate = train_test_split(  
    df,  
    test_size=0.30, # 30% reserved for test + validation  
    random_state=42,  
    stratify=df['target'])  
  
validate, test = train_test_split(  
    test_validate,  
    test_size=0.50, # split the 30% into two halves = 15% each  
    random_state=42,  
    stratify=test_validate['target'])  
  
print("Train:", train.shape)  
print("Validate:", validate.shape)  
print("Test:", test.shape)  
  
print("\nTrain distribution:\n", train['target'].value_counts(normalize=True))  
print("\nValidate distribution:\n", validate['target'].value_counts(normalize=True))  
print("\nTest distribution:\n", test['target'].value_counts(normalize=True))
```

```
Train: (2412, 86)
Validate: (517, 86)
Test: (518, 86)

Train distribution:
target
0.0    0.800995
1.0    0.199005
Name: proportion, dtype: float64

Validate distribution:
target
0.0    0.800774
1.0    0.199226
Name: proportion, dtype: float64

Test distribution:
target
0.0    0.801158
1.0    0.198842
Name: proportion, dtype: float64
```

2. Use xgboost estimator to build a classification model.

```
In [12]: bucket='c182567a4701747l12448301t1w754121081532-labbucket-z0zjck4jphe5'
```

```
In [13]: prefix='lab3'
```

```
train_file='new_delay_train.csv'
test_file='new_delay_test.csv'
validate_file='new_delay_validate.csv'

s3_resource = boto3.Session().resource('s3')
def upload_s3_csv(filename, folder, dataframe):
    csv_buffer = io.StringIO()
    dataframe.to_csv(csv_buffer, header=False, index=False)
    s3_resource.Bucket(bucket).Object(os.path.join(prefix, folder, filename)).put(Body=csv_buffer.getvalue())
```

```
In [14]: upload_s3_csv(train_file, 'train', train)
```

```
upload_s3_csv(test_file, 'test', test)
```

```
upload_s3_csv(validate_file, 'validate', validate)
```

```
In [15]: for obj in s3_resource.Bucket(bucket).objects.filter(Prefix=prefix):
    print(obj.key)
```

```
lab3/test/new_delay_test.csv
```

```
lab3/train/new_delay_train.csv
```

```
lab3/validate/new_delay_validate.csv
```

```
In [16]: container = retrieve('xgboost',boto3.Session().region_name,'1.0-1')
```

```
In [17]: hyperparams={"num_round":42,
                     "eval_metric": "auc",
                     "objective": "binary:logistic"}
```

```
In [18]: s3_output_location="s3://{}:{}//output/".format(bucket,prefix)
xgb_model=sagemaker.estimator.Estimator(container,
                                         sagemaker.get_execution_role(),
                                         instance_count=1,
                                         instance_type='ml.m4.xlarge',
                                         output_path=s3_output_location,
                                         hyperparameters=hyperparams,
                                         sagemaker_session=sagemaker.Session())
```

```
In [19]: train_channel = sagemaker.inputs.TrainingInput(
    "s3://{}:{}//train/".format(bucket,prefix,train_file),
    content_type='text/csv')

validate_channel = sagemaker.inputs.TrainingInput(
    "s3://{}:{}//validate/".format(bucket,prefix,validate_file),
    content_type='text/csv')

data_channels = {'train': train_channel, 'validation': validate_channel}

xgb_model.fit(inputs=data_channels, logs=False)
```

INFO:sagemaker:Creating training-job with name: sagemaker-xgboost-2025-11-03-12-04-20-403

2025-11-03 12:04:20 Starting - Starting the training job.....
2025-11-03 12:05:04 Starting - Preparing the instances for training.....
2025-11-03 12:05:39 Downloading - Downloading input data.....
2025-11-03 12:06:09 Downloading - Downloading the training image.....
2025-11-03 12:07:15 Training - Training image download completed. Training in progress....
2025-11-03 12:07:36 Uploading - Uploading generated training model...
2025-11-03 12:07:54 Completed - Training job completed

3. Host the model on another instance

```
In [20]: # Deploy the trained Linear Learner model
xgb_predictor = xgb_model.deploy(
    initial_instance_count=1,
    instance_type='ml.m4.xlarge',
)

print(f"Model deployed successfully!")
```

```
INFO:sagemaker:Creating model with name: sagemaker-xgboost-2025-11-03-12-08-46-466
INFO:sagemaker:Creating endpoint-config with name sagemaker-xgboost-2025-11-03-12-08-46-466
INFO:sagemaker:Creating endpoint with name sagemaker-xgboost-2025-11-03-12-08-46-466
```

```
-----!Model deployed successfully!
```

4. Perform batch transform to evaluate the model on testing data

```
In [21]: # Prepare test data (features only, no target)
batch_X = test.iloc[:, 1:] # Exclude target column (first column)

# Upload batch input to S3
batch_X_file = 'batch-in.csv'
upload_s3_csv(batch_X_file, 'batch-in', batch_X)

# Define S3 paths
batch_output = "s3://{}/{}/batch-out/".format(bucket, prefix)
batch_input = "s3://{}/{}/batch-in/{}".format(bucket, prefix, batch_X_file)

print("Starting batch transform...")

xgb_transformer = xgb_model.transformer(instance_count=1,
                                         instance_type='ml.m4.xlarge',
                                         strategy='MultiRecord',
                                         assemble_with='Line',
                                         output_path=batch_output)

xgb_transformer.transform(data=batch_input,
                         data_type='S3Prefix',
                         content_type='text/csv',
                         split_type='Line')
xgb_transformer.wait()
print(" Batch transform completed!")
```

```
INFO:sagemaker:Creating model with name: sagemaker-xgboost-2025-11-03-12-13-25-057
```

```
Starting batch transform...
```

```
INFO:sagemaker:Creating transform job with name: sagemaker-xgboost-2025-11-03-12-13-25-638
```

```
.....
```

```
... Batch transform completed!
```

```
In [22]: s3 = boto3.client('s3')
obj = s3.get_object(Bucket=bucket, Key=f'{prefix}/batch-out/{batch_X_file}.out')

# XGBoost outputs plain floats (probabilities)
predictions_raw = obj['Body'].read().decode('utf-8').strip().split('\n')

# Parse scores
predicted_scores = [float(line.strip()) for line in predictions_raw]

print(f"  Total predictions: {len(predicted_scores)}")
print(f"  Score range: [{min(predicted_scores):.4f}, {max(predicted_scores):.4f}]")
print(f"  First 5 scores: {predicted_scores[:5]}")

# Apply threshold
threshold = 0.5
predicted_labels = [1 if score > threshold else 0 for score in predicted_scores]

# Create DataFrame
target_predicted = pd.DataFrame({
    'score': predicted_scores,
    'predicted': predicted_labels
})

print(f"\n  Prediction distribution (threshold={threshold}):")
print(f"  - Predicted Delayed (1): {sum(predicted_labels)}, ({100*sum(predicted_labels)}/len(predicted_labels)}%)")
print(f"  - Predicted No Delay (0): {len(predicted_labels)-sum(predicted_labels)}, ({100*(len(predicted_labels)-sum(predicted_labels))}/len(predicted_labels)}%)")
```

```
Total predictions: 518
Score range: [0.0031, 0.9506]
First 5 scores: [0.09757790714502335, 0.08573710173368454, 0.27737271785736084, 0.36310312151908875, 0.387405276298523]
```

```
Prediction distribution (threshold=0.5):
- Predicted Delayed (1): 49 (9.5%)
- Predicted No Delay (0): 469 (90.5%)
```

```
In [23]: def binary_convert(x):
    threshold = 0.3
    if x > threshold:
        return 1
    else:
        return 0

target_predicted_binary = target_predicted['score'].apply(binary_convert)

print("Prediction DataFrame structure:")
print(target_predicted.head(10))
print("\nColumns:", target_predicted.columns.tolist())
print("\nData types:")
print(target_predicted.dtypes)
```

Prediction DataFrame structure:

	score	predicted
0	0.097578	0
1	0.085737	0
2	0.277373	0
3	0.363103	0
4	0.338741	0
5	0.451829	0
6	0.727647	1
7	0.041521	0
8	0.033220	0
9	0.228706	0

Columns: ['score', 'predicted']

Data types:

score	float64
predicted	int64
dtype:	object

```
In [27]: # CALCULATE CONFUSION MATRIX AND METRICS FIRST
from sklearn.metrics import (confusion_matrix, roc_curve, auc,
                             accuracy_score, precision_score, recall_score,
                             f1_score, roc_auc_score, classification_report)
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

# Get true labels and predictions
y_true = test['target'].values
y_pred = np.array(predicted_labels)
y_scores = np.array(predicted_scores)

print("XGBoost - PERFORMANCE METRICS (combined_csv_v2.csv)")

# Calculate all metrics
accuracy = accuracy_score(y_true, y_pred)
precision = precision_score(y_true, y_pred)
recall = recall_score(y_true, y_pred)
f1 = f1_score(y_true, y_pred)
roc_auc = roc_auc_score(y_true, y_scores)

print(f"\n1. Accuracy: {accuracy:.4f}")
print(f"2. Precision: {precision:.4f}")
print(f"3. Recall: {recall:.4f}")
print(f"4. F1-Score: {f1:.4f}")
print(f"5. ROC-AUC: {roc_auc:.4f}")

# Calculate confusion matrix
cm = confusion_matrix(y_true, y_pred)

print("Confusion Matrix:")
print(cm)
print(f"\nTrue Negatives (TN): {cm[0,0]},")
print(f"False Positives (FP): {cm[0,1]},")
print(f"False Negatives (FN): {cm[1,0]},")
print(f"True Positives (TP): {cm[1,1]}")
```

XGBoost – PERFORMANCE METRICS (combined_csv_v2.csv)

1. Accuracy: 0.7992
2. Precision: 0.4898
3. Recall: 0.2330
4. F1-Score: 0.3158
5. ROC-AUC: 0.7006

Confusion Matrix:

```
[[390  25]
 [ 79  24]]
```

True Negatives (TN): 390

False Positives (FP): 25

False Negatives (FN): 79

True Positives (TP): 24

```
In [26]: import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import roc_curve, auc

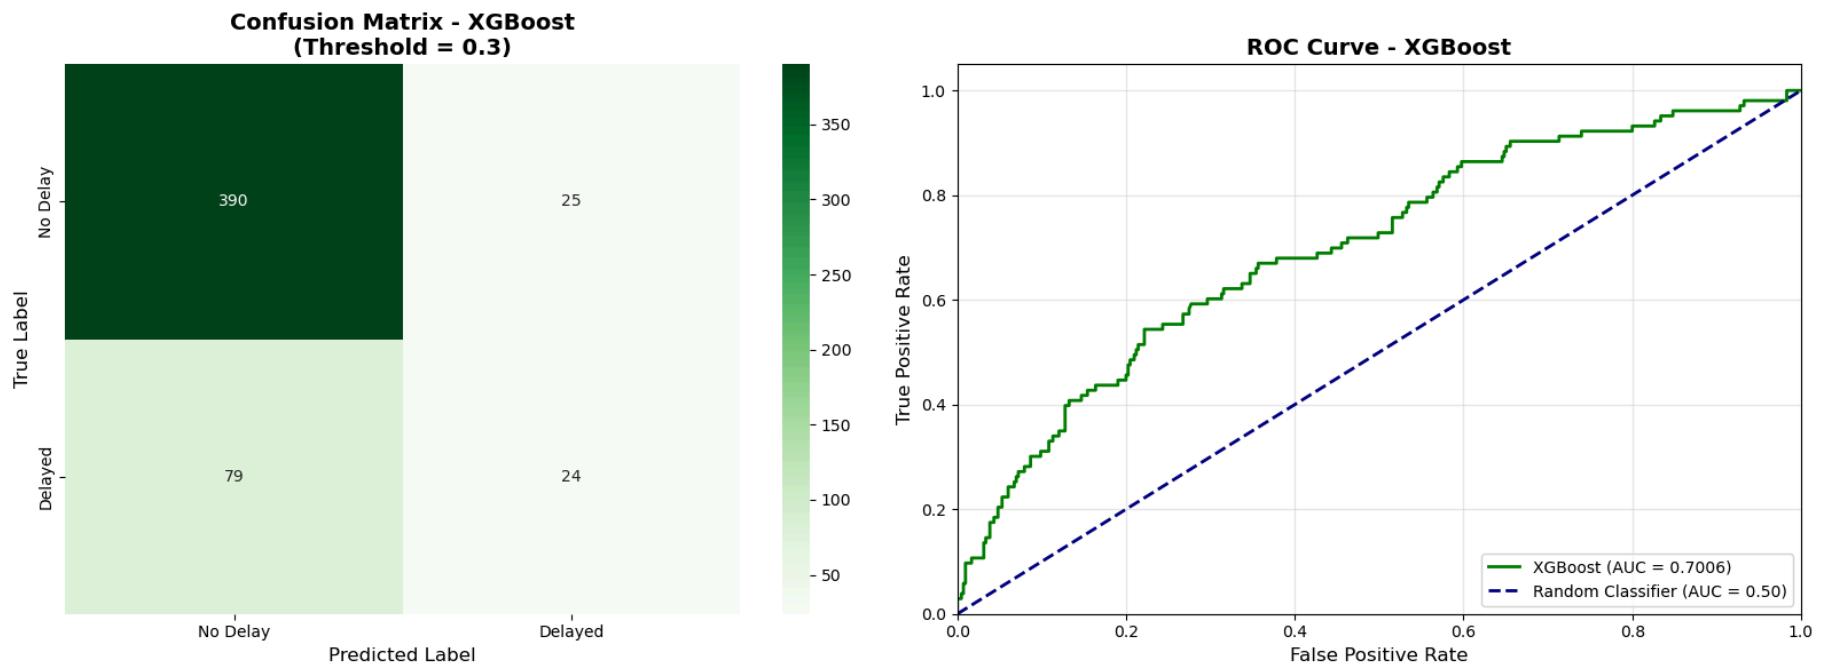
# Create figure with subplots
fig, axes = plt.subplots(1, 2, figsize=(16, 6))

# 1. Confusion Matrix Heatmap
sns.heatmap(cm, annot=True, fmt='d', cmap='Greens', ax=axes[0],
            xticklabels=['No Delay', 'Delayed'],
            yticklabels=['No Delay', 'Delayed'])
axes[0].set_title('Confusion Matrix - XGBoost\n(Threshold = 0.3)', fontsize=14, fontweight='bold')
axes[0].set_ylabel('True Label', fontsize=12)
axes[0].set_xlabel('Predicted Label', fontsize=12)

# 2. ROC Curve
fpr, tpr, thresholds = roc_curve(y_true, y_scores)
roc_auc_calc = auc(fpr, tpr)

axes[1].plot(fpr, tpr, color='green', lw=2,
             label=f'XGBoost (AUC = {roc_auc_calc:.4f})')
axes[1].plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--',
             label='Random Classifier (AUC = 0.50)')
axes[1].set_xlim([0.0, 1.0])
axes[1].set_ylim([0.0, 1.05])
axes[1].set_xlabel('False Positive Rate', fontsize=12)
axes[1].set_ylabel('True Positive Rate', fontsize=12)
axes[1].set_title('ROC Curve - XGBoost', fontsize=14, fontweight='bold')
axes[1].legend(loc="lower right", fontsize=10)
axes[1].grid(alpha=0.3)

plt.tight_layout()
plt.show()
```



write down your observation on the difference between the performance of using the simple and ensemble models. Note: You are required to perform the above steps on the two combined datasets separatey.

The comparison between Linear Learner and XGBoost reveals that the ensemble model significantly outperforms the simple model. XGBoost achieved a ROC-AUC of 0.7006 compared to Linear Learner's 0.5765, demonstrating 21% better discriminative ability. XGBoost also shows superior accuracy (79.92% vs 68.56%) and precision (48.98% vs 27.11%), meaning it makes far fewer false predictions when identifying delays.

However, both models struggle with recall, which is the most critical metric for this business problem. Linear Learner catches only 29.82% of actual delays, while XGBoost performs even worse at 23.30%, missing approximately 70-77% of delayed flights. This means most customers wouldn't receive advance warnings, defeating the system's primary purpose. The poor recall occurs because both models use a conservative 0.5 probability threshold that prioritizes avoiding false alarms over catching delays.

In conclusion, XGBoost is clearly superior. If I have more time, I will focus on tuning the classification threshold and optimizing recall to better capture actual flight delays.

In []: